

DTS202TC Foundation of Parallel Computing

DTS202TC Foundation of Parallel Computing

I. Lab 2 (30 marks)

1.1 Task1

1.1.1 Code

1.1.2 Screenshot

1.2 Task2

1.2.1 Code

1.2.2 Screenshot

1.3 Task3

1.3.1 Code

1.3.2 Screenshot

1.3.3 Analysis the performance

II. Lab 3 (35 Marks)

2.1 Profiling

2.1.1 Code

2.1.2 Screenshot

2.1.3 Analysis.txt

2.2 Pthreads Hello World

2.2.1 Code

2.2.2 Screenshot

2.3 Estimate π with Pthreads

2.3.1 Code

2.3.2 Screenshot

III. Lab 4 (35 Marks)

3.1 OpenMP Hello World

3.1.1 Code

3.1.2 Screenshot

3.2 Estimate π with OpenMP

3.2.1 Code

3.2.2 Screenshot

3.3 Comparison

3.3.1 Comparison under different threads

3.3.1.1 Elapsed Time

3.3.1.2 Speedup

3.3.1.3 Efficiency

3.3.2 Comparison under different problem size

I. Lab 2 (30 marks)

1.1 Task1

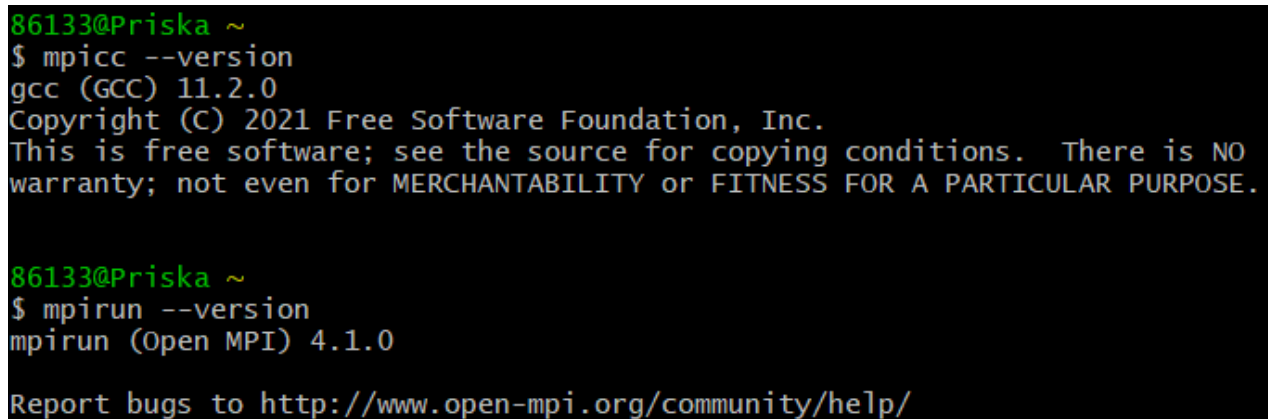
The author's computer system is Windows10, so using Cygwin to compile `.c` files. OpenMPI and related Package have been successfully installed on PC, and `.c` files with mpi header files can be successfully compiled on the computer.

Prove that the installation is successful by using a line command and compiling a test `.c` file

1.1.1 Code

```
1 mpicc --version
2 mpirun --version
```

1.1.2 Screenshot



```
86133@Priska ~
$ mpicc --version
gcc (GCC) 11.2.0
Copyright (C) 2021 Free Software Foundation, Inc.
This is free software; see the source for copying conditions.  There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.

86133@Priska ~
$ mpirun --version
mpirun (Open MPI) 4.1.0

Report bugs to http://www.open-mpi.org/community/help/
```

1.2 Task2

1.2.1 Code

Shell

```
1 mpicc -g -Wall -o mpi_hello mpi_hello.c
2 mpiexec -n 4 ./mpi_hello
3 mpiexec -n 1 ./mpi_hello
4 mpiexec -n 2 ./mpi_hello
```

.c

```
1 #include <stdio.h>
2 #include <string.h>
3 #include <mpi.h>
4
5 const int MAX_STRING = 100;
6
7 int main(void){
8     char greeting[MAX_STRING];
9     int comm_sz;
10    int my_rank;
11
12    MPI_Init(NULL, NULL);
13    MPI_Comm_size(MPI_COMM_WORLD, &comm_sz);
14    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
15
16    if(my_rank != 0){
17        sprintf(greeting, "Greetings from process %d of %d",
18                my_rank, comm_sz);
19
20        MPI_Send(greeting, strlen(greeting)+1, MPI_CHAR, 0, 0, MPI_COMM_WORLD);
21    } else{
22        printf("Greeting from process %d of %d! \n", my_rank,
23              comm_sz);
24        for (int q = 1; q < comm_sz; q++) {
```

```
24     MPI_Recv(greeting, MAX_STRING, MPI_CHAR, q, 0, MPI_COMM_WORLD, MPI_ST
ATUS_IGNORE);
25         printf("%s\n", greeting);
26     }
27 }
28
29 MPI_Finalize();
30 return 0;
31 }
```

1.2.2 Screenshot

```
86133@Priska /cygdrive/d/dts202tc/MPI_Lab/Lab2_1
$ ls
CMakeLists.txt  cmake-build-debug  mpi_hello.c  serial_pi.c  task1.exe

86133@Priska /cygdrive/d/dts202tc/MPI_Lab/Lab2_1
$ mpicc -g -Wall -o mpi_hello mpi_hello.c

86133@Priska /cygdrive/d/dts202tc/MPI_Lab/Lab2_1
$ mpiexec -n 4 ./mpi_hello
Greeting from process 0 of 4!
Greetings from process 1 of 4
Greetings from process 2 of 4
Greetings from process 3 of 4
```

```
86133@Priska /cygdrive/d/dts202tc/MPI_Lab/Lab2_1
$ mpiexec -n 1 ./mpi_hello
Greeting from process 0 of 1!
```

```
86133@Priska /cygdrive/d/dts202tc/MPI_Lab/Lab2_1
$ mpiexec -n 2 ./mpi_hello
Greeting from process 0 of 2!
Greetings from process 1 of 2
```

1.3 Task3

1.3.1 Code

```
1 gcc -g -Wall -o serial_pi serial_pi.c
2 ./serial_pi
3
4 mpicc -g -Wall -o parallel_pi parallel_pi.c
5 mpiexec -n 1 ./parallel_pi.c
6 mpiexec -n 2 ./parallel_pi.c
7 mpiexec -n 3 ./parallel_pi.c
8 mpiexec -n 4 ./parallel_pi.c
```

- serial_pi.c

```
1 #include <stdio.h>
2 #include <math.h>
3 #include <stdlib.h>
4 #include <sys/time.h>
5
6 #define GET_TIME(now){ \
7     struct timeval t;    \
8     gettimeofday(&t,NULL);\
9     now = t.tv_sec+t.tv_sec/1000000.0;\
10 }
11
12 double serial_pi(long long n);
13 int main(int argc, char** argv) {
14     double start, finish;
15     GET_TIME(start);
16     double estimate_of_pi = serial_pi(1000000000);
17     printf("\n Estimated of pi: %1.10f.\n", estimate_of_pi);
18     GET_TIME(finish);
19
20     printf("\nActual value of pi: %1.10f.\n\n", atan(1)*4);
21     printf("The elapsed time is %e seconds\n", finish-start);
22 }
23
24 double serial_pi(long long n){
25     double sum = 0.0;
```

```

26     long long i;
27     double factor = 1.0;
28
29     for (i = 0; i < n; i++, factor=-factor) {
30         sum += factor/(2*i+1);
31     }
32
33     return 4.0*sum;
34 }

```

- parallel_pi.c

```

1  /*File:
2   *  parallel_pi.c
3   *
4   *Purpose:
5   * Using MPI_Send and MPI_Recv to estimate_pi
6   *
7   *Compile:
8   * mpicc -g -Wall -o parallel_pi parallel_pi.c
9   *Usage:
10  * mpiexec -n<number of processes> ./parallel_pi
11  *
12  *Input:
13  *  None
14  *Output:
15  * n processes estimates the value of the PI
16  *
17  * Algorithm:
18  *  1. Each process calculates its interval of integration pi
19  *  2. Each process != 0 sends its integral to 0 process
20  *  3. Process 0 sums the calculations received from the
    individual
21  *      processes and prints the result
22  * */
23  #include <stdio.h>
24  #include <mpi.h>
25  #include <math.h>
26  #include <sys/time.h>
27
28  #define GET_TIME(now){ \
29      struct timeval t;    \

```

```

30     gettimeofday(&t,NULL);\
31     now = t.tv_sec+t.tv_sec/1000000.0;\
32 }
33 double serial_pi(long long a,long long b);/*Calculate local
integral*/
34
35 int main(void) {
36     int my_rank,comm_sz,source;/*My process rank, Number of
processes*/
37     long long n=1000000000,local_n,local_a,local_b;
38     double local_est, total_est;
39     double local_start,local_finish,local_elapsed,elapsed;
40
41     MPI_Init(NULL,NULL);
42     MPI_Comm_rank(MPI_COMM_WORLD,&my_rank);
43     MPI_Comm_size(MPI_COMM_WORLD,&comm_sz);
44
45     local_n = n/comm_sz;
46     local_a = 0+my_rank*local_n;
47     local_b = local_a+local_n;
48
49
50     MPI_Barrier(MPI_COMM_WORLD);
51     local_start=MPI_Wtime();
52     local_est = serial_pi(local_a,local_b);
53
54     if (my_rank!=0){
55         MPI_Send(&local_est,1,MPI_DOUBLE,0,0,MPI_COMM_WORLD);
56
57     } else{
58         total_est=local_est;
59         for (source = 1; source < comm_sz; source++) {
60
61             MPI_Recv(&local_est,1,MPI_DOUBLE,source,0,MPI_COMM_WORLD,MPI_ST
ATUS_IGNORE);
62             total_est+=local_est;
63         }
64     }
65     local_finish = MPI_Wtime();
66     local_elapsed = local_finish-local_start;
67
68     MPI_Reduce(&local_elapsed,&elapsed,1,MPI_DOUBLE,MPI_MAX,0,MPI_C
OMM_WORLD);

```

```

67
68     if (my_rank==0){
69         printf("\n Estimated of pi: %1.10f.\n",total_est);
70         printf("\nAcutal value of pi: %1.10f.\n\n",
atan(1)*4);
71     }
72     if(my_rank==0) {
73         printf(" The elapsed time is %e seconds\n", elapsed);
74     }
75     MPI_Finalize();
76
77     return 0;
78 }
79
80 double serial_pi(long long a,long long b){
81     double sum = 0.0;
82     double factor;
83     if (a/2 == 0){
84         factor = 1.0;
85     } else{
86         factor = -1.0;
87     }
88
89     for (; a < b; a++,factor=-factor) {
90         sum += factor/(2*a+1);
91     }
92
93     return 4.0*sum;
94 }

```

1.3.2 Screenshot

```

86133@Priska /cygdrive/d/dts202tc/MPI_Lab/Lab2_4
$ gcc -g -Wall -o serial_pi serial_pi.c

86133@Priska /cygdrive/d/dts202tc/MPI_Lab/Lab2_4
$ ./serial_pi

Estimated of pi: 3.1415926526.

Acutal value of pi: 3.1415926536.

The elapsed time is 2.000002e+00 seconds

```



```
86133@Priska /cygdrive/d/dts202tc/mpi_lab/lab2_4
$ mpiexec -n 1 ./parallel_pi
Estimated of pi: 3.1415926526.
Acutal value of pi: 3.1415926536.
The elapsed time is 2.377387e+00 seconds
```

```
86133@Priska /cygdrive/d/dts202tc/mpi_lab/lab2_4
$ mpiexec -n 2 ./parallel_pi
Estimated of pi: 3.1415926506.
Acutal value of pi: 3.1415926536.
The elapsed time is 1.224703e+00 seconds
```

```
86133@Priska /cygdrive/d/dts202tc/mpi_lab/lab2_4
$ mpiexec -n 3 ./parallel_pi
Estimated of pi: 3.1415926496.
Acutal value of pi: 3.1415926536.
The elapsed time is 9.177730e-01 seconds
```

```
86133@Priska /cygdrive/d/dts202tc/mpi_lab/lab2_4
$ mpiexec -n 4 ./parallel_pi
Estimated of pi: 3.1415926466.
Acutal value of pi: 3.1415926536.
The elapsed time is 8.175870e-01 seconds
```

1.3.3 Analysis the performance

	1 PROCESSOR	2 PROCESSOR	3 PROCESSOR	4 PROCESSOR
Serial_pi	2.00 s			
Parallel_pi	2.37 s	1.22 s	0.92 s	0.82 s

Generally, the parallel program of 1 process should have the same time as the serial program, but there is a difference of time. This is due to the different time calculation functions used by the two programs. The parallel program uses the following time function to time the mpi program

```

1  double local_start,local_finish,local_elapsed,elapsed;
2  ...
3  MPI_Barrier(comm);
4  local_start = MPI_Wtime();
5  /*Code to be timed*/
6  ...
7  local_finish = MPI_Wtime();
8  local_elapsed = local_finish - local_start;
9  MPI_Reduce(&local_elapsed,&elapsed,1,MPI_DOUBLE,MPI_MAX,0,comm)
   ;
10
11  if(my_rank==0){
12      printf("Elapsed time = %e seconds\n",elapsed);
13  }

```

Such a timing function causes time overhead, so it is slightly different from serial time.

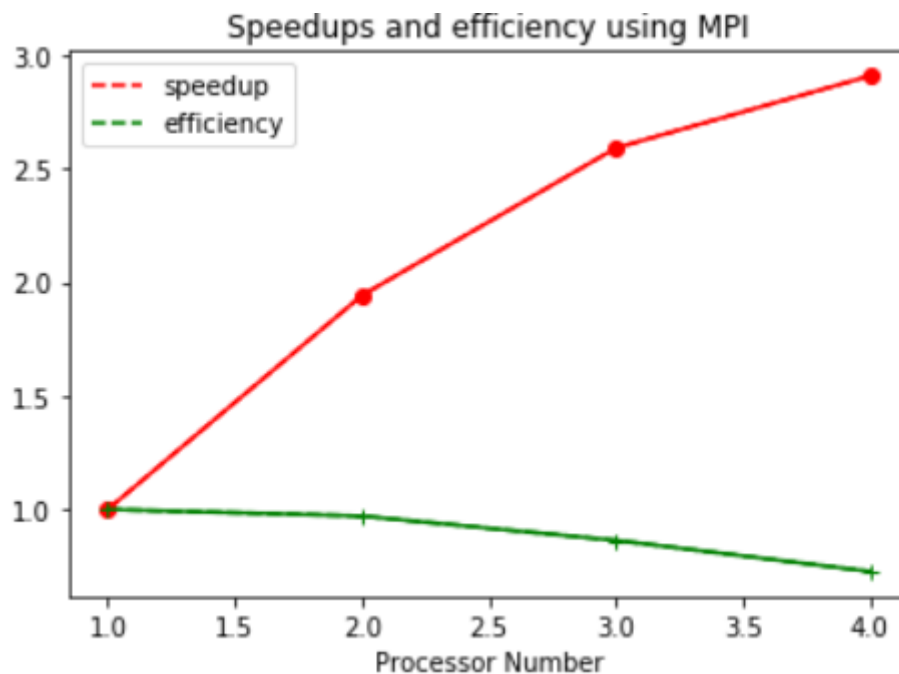
MPI Parallel Algorithm:

- Each process calculates its interval of integration π
- Each process $\neq 0$ sends its integral to 0 process
- Process 0 sums the calculations received from the individual processes and prints the result

Such a parallel algorithm significantly shortens the time of serial programs.

Then compare the performance of different processor by analyzing speedup and efficiency

	1 PROCESSOR	2 PROCESSORS	3 PROCESSORS	4 PROCESSORS
speedup	1	1.94	2.59	2.91
efficiency	1	0.97	0.86	0.73



When the size of the problem remains the same, the speedup increases as the number of processes increases, because the more processes, the less time the program takes to execute. Ideally, speedup should be linear speedup, that is, the number of processes. And there will be an upper limit according to Amdahl's Law, speedup, but due to the limit on the number of machine processor, it is not possible to try more processes to see when the limit is reached.

However, it is found that the efficiency is decreasing, which is that when the Processor increases, although the time decreases, the relative communication overhead also increases, that is, the more processor, the greater the communication overhead, the lower the efficiency.

II. Lab 3 (35 Marks)

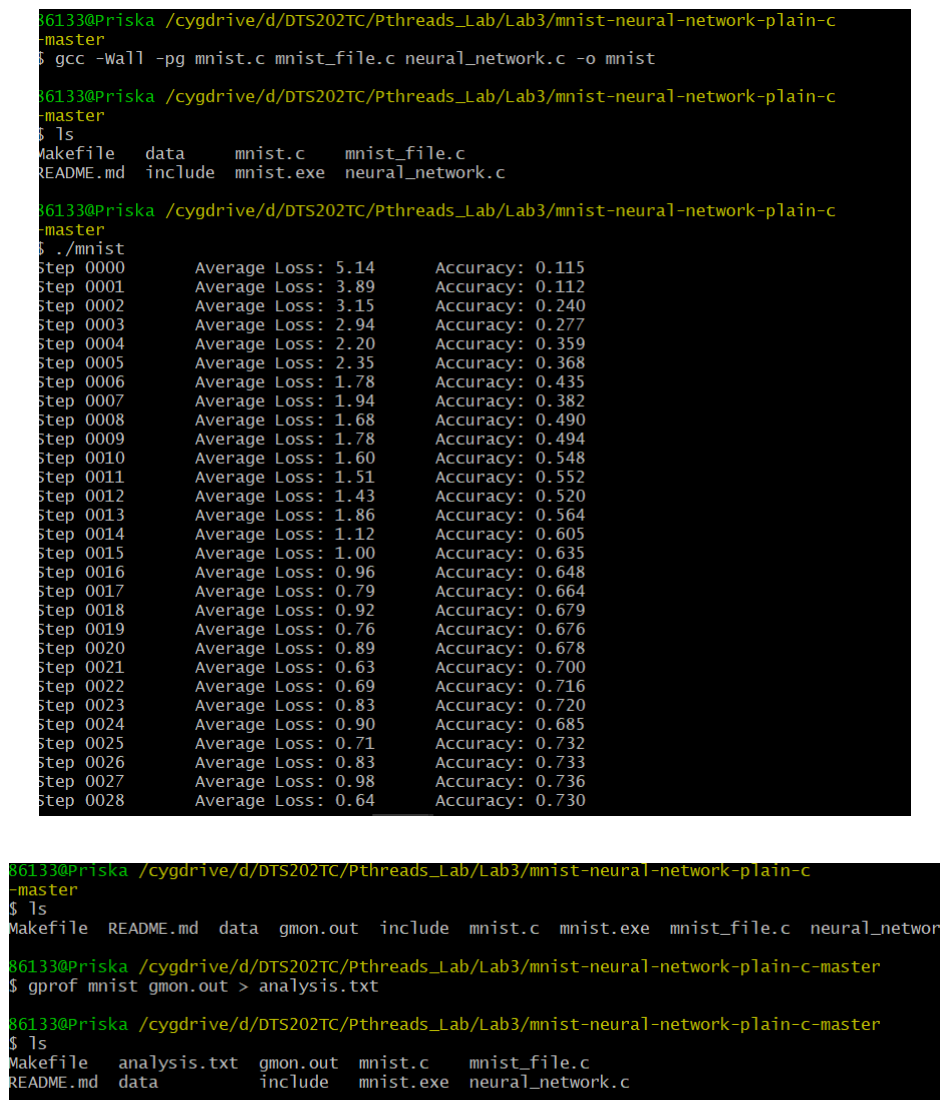
2.1 Profiling

2.1.1 Code

- On Cygwin

```
1 gcc -Wall -pg mnist.c mnist_file.c neural_network.c -o mnist
2 ./mnist
3 gprof mnist gmon.out > analysis.txt
```

2.1.2 Screenshot



The first screenshot shows the compilation of the program using gcc with profiling flags (-Wall -pg) and the execution of the resulting binary (mnist). The output displays a table of Average Loss and Accuracy for each step from 0000 to 0028.

Step	Average Loss	Accuracy
0000	5.14	0.115
0001	3.89	0.112
0002	3.15	0.240
0003	2.94	0.277
0004	2.20	0.359
0005	2.35	0.368
0006	1.78	0.435
0007	1.94	0.382
0008	1.68	0.490
0009	1.78	0.494
0010	1.60	0.548
0011	1.51	0.552
0012	1.43	0.520
0013	1.86	0.564
0014	1.12	0.605
0015	1.00	0.635
0016	0.96	0.648
0017	0.79	0.664
0018	0.92	0.679
0019	0.76	0.676
0020	0.89	0.678
0021	0.63	0.700
0022	0.69	0.716
0023	0.83	0.720
0024	0.90	0.685
0025	0.71	0.732
0026	0.83	0.733
0027	0.98	0.736
0028	0.64	0.730

The second screenshot shows the execution of gprof to generate a profiling report (analysis.txt) and the subsequent listing of files in the directory.

2.1.3 Analysis.txt

Flat profile:

Each sample counts as 0.01 seconds.

%	cumulative	self		self	total	
time	seconds	seconds	calls	ms/call	ms/call	name
97.87	725.00	725.00	10100000	0.07	0.07	neural_network_hypothesis
1.06	732.88	7.88	100000	0.08	0.15	neural_network_gradient_update
0.43	736.09	3.21	10100000	0.00	0.00	neural_network_softmax
0.28	738.19	2.10	1000	2.10	723.10	calculate_accuracy
0.23	739.90	1.71				_mcount_private
0.08	740.52	0.62				__fentry__
0.02	740.68	0.16	1000	0.16	15.25	neural_network_training_step
0.01	740.78	0.10	1	100.00	100.00	neural_network_random_weights
0.00	740.78	0.00	1000	0.00	0.00	mnist_batch
0.00	740.78	0.00	12	0.00	0.00	map_uint32
0.00	740.78	0.00	2	0.00	0.00	get_images
0.00	740.78	0.00	2	0.00	0.00	get_labels
0.00	740.78	0.00	2	0.00	0.00	mnist_free_dataset
0.00	740.78	0.00	2	0.00	0.00	mnist_get_dataset

index	% time	self	children	called	name
[1]	99.7	0.00	738.45		<spontaneous>
		2.10	721.00	1000/1000	main [1]
		0.16	15.09	1000/1000	calculate_accuracy [3]
		0.10	0.00	1/1	neural_network_training_step [4]
		0.00	0.00	1000/1000	neural_network_random_weights [9]
		0.00	0.00	2/2	mnist_batch [10]
		0.00	0.00	2/2	mnist_get_dataset [15]
		0.00	0.00	2/2	mnist_free_dataset [14]
		7.18	0.03	100000/10100000	neural_network_gradient_update [5]
		717.83	3.17	10000000/10100000	calculate_accuracy [3]
[2]	98.3	725.00	3.21	10100000	neural_network_hypothesis [2]
		3.21	0.00	10100000/10100000	neural_network_softmax [6]
		2.10	721.00	1000/1000	main [1]
[3]	97.6	2.10	721.00	1000	calculate_accuracy [3]
		717.83	3.17	10000000/10100000	neural_network_hypothesis [2]
		0.16	15.09	1000/1000	main [1]
[4]	2.1	0.16	15.09	1000	neural_network_training_step [4]
		7.88	7.21	100000/100000	neural_network_gradient_update [5]
		7.88	7.21	100000/100000	neural_network_training_step [4]
[5]	2.0	7.88	7.21	100000	neural_network_gradient_update [5]
		7.18	0.03	100000/10100000	neural_network_hypothesis [2]
		3.21	0.00	10100000/10100000	neural_network_hypothesis [2]
[6]	0.4	3.21	0.00	10100000	neural_network_softmax [6]
					<spontaneous>
[7]	0.2	1.71	0.00		_mcount_private [7]
					<spontaneous>
[8]	0.1	0.62	0.00		__fentry__ [8]
		0.10	0.00	1/1	main [1]
[9]	0.0	0.10	0.00	1	neural_network_random_weights [9]
		0.00	0.00	1000/1000	main [1]
[10]	0.0	0.00	0.00	1000	mnist_batch [10]

		0.00	0.00	4/12	get_labels [13]
		0.00	0.00	8/12	get_images [12]
[11]	0.0	0.00	0.00	12	map_uint32 [11]
		0.00	0.00	2/2	mnist_get_dataset [15]
[12]	0.0	0.00	0.00	2	get_images [12]
		0.00	0.00	8/12	map_uint32 [11]
		0.00	0.00	2/2	mnist_get_dataset [15]
[13]	0.0	0.00	0.00	2	get_labels [13]
		0.00	0.00	4/12	map_uint32 [11]
		0.00	0.00	2/2	main [1]
[14]	0.0	0.00	0.00	2	mnist_free_dataset [14]
		0.00	0.00	2/2	main [1]
[15]	0.0	0.00	0.00	2	mnist_get_dataset [15]
		0.00	0.00	2/2	get_images [12]
		0.00	0.00	2/2	get_labels [13]

The profiling analysis of mnist-neural-network-plain-c-master is carried out by using GPROF tool. As can be seen from the above figure, `neural_network_hypothesis` takes it as the bottleneck of this program as the function with the longest proportion. Later improvements and upgrades take this function as the key object of study.

2.2 Pthreads Hello World

2.2.1 Code

- On Cygwin

```
1 gcc -g -Wall -o pth_hello pth_hello.c -lpthread
2 ./pth_hello 4
```

- pth_hello.c

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <pthread.h>
4
5 /*Global variable: accessible to all threads*/
6 int thread_count;
7 void* Hello(void* rank);
8
9 int main(int argc, char* argv[]) {
10     long thread;    /*Using long in case of a 64-bit system*/
11     pthread_t* thread_handles;
12
13     thread_count = strtol(argv[1], NULL, 10);
14
15     thread_handles = malloc(thread_count*sizeof(pthread_t));
16
17     for (thread = 0; thread < thread_count; thread++) {
18         pthread_create(&thread_handles[thread], NULL,
19                     Hello, (void*)thread);
20     }
21
22     printf("Hello from the main thread\n");
23
24     for (thread = 0; thread < thread_count; thread++) {
25         pthread_join(thread_handles[thread], NULL);
26     }
27
28     free(thread_handles);
29     return 0;
```

```
30 }
31
32 void* Hello(void* rank){
33     long my_rank = (long) rank;
34
35     printf("Hello from thread %ld of %d\n",
36           my_rank, thread_count);
37
38     return NULL;
39 }
```

2.2.2 Screenshot

```
86133@Priska /cygdrive/d/dts202tc/Pthreads_Lec/pth_hello
$ gcc -g -Wall -o pth_hello pth_hello.c -lpthread

86133@Priska /cygdrive/d/dts202tc/Pthreads_Lec/pth_hello
$ ./pth_hello 4
Hello from thread 0 of 4
Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 3 of 4
Hello from the main thread
```


2.3 Estimate π with Pthreads

The program originally used `random()` for random numbers in the `Threadsum` function, but this function does not improve performance through threads, and should use the global state shared between threads, there is no way to guarantee thread safe. Finally, the program chooses `rand_r()` function to improve performance.

2.3.1 Code

- shell

```
1 gcc -g -Wall -o pth_pi pth_pi.c -lm -lpthread
2 ./pth_pi 10000000 1
3 ./pth_pi 10000000 2
4 ./pth_pi 10000000 4
5 ./pth_pi 10000000 6
6 ./pth_pi 10000000 8
7 ./pth_pi 10000000 10
8 ./pth_pi 10000000 12
9 ./pth_pi 10000000 14
10 ./pth_pi 10000000 16
```

- `pth_pi.c`

```
1 /*File: pth_pi.c
2  * Purpose: Estimate pi using Monte Carlo Method
3  *         (number in circle)/(total number of tosses) = pi/4
4  *         This version used a mutex to protect the critical
5  *         section
6  *
7  * Compile: gcc -g -Wall -o pth_pi pth_pi.c -lm -lpthread
8  * Run:     ./pth_pi <total_points><number of threads>
9  *         total_points is the randomly toss total
10 *         should be evenly divisible by the number of
11 *         threads
12 *
13 *Input:     None
14 *Output:    The estimate of pi using Monte Carlo on multiple
15 *         threads,
```

```

13      *           Also elapsed times for the multithreaded and
           singlethreaded
14      *           computations
15      */
16
17 #include <pthread.h>
18 #include <stdio.h>
19 #include <stdlib.h>
20 #include <time.h>
21
22 #ifndef _TIMER_H_
23 #define _TIMER_H_
24 #include <sys/time.h>
25 #include <stdlib.h>
26 /* The argument now should be a double (not a pointer to a
           double) */
27 #define GET_TIME(now) { \
28     struct timeval t; \
29     gettimeofday(&t, NULL); \
30     now = t.tv_sec + t.tv_usec/1000000.0; \
31 }
32 #endif
33
34 /*Global variables: accessible to all threads*/
35 const int MAX_THREADS = 1024; /*Threads maximum*/
36 long long in_circle = 0;      /*Each thread has in_circle
           points*/
37 long long points_per_thread; /*Each thread count points
           number*/
38 long long total_points;      /*total points from user input*/
39 long long thread_count;      /*total threads from user input*/
40 pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER; /*Initialize
           mutex*/
41
42 /*Only executed by main thread*/
43 void* Thread_sum();
44 void Get_args(int argc, char*argv[]);
45 void Usage(char* prog_name);
46
47 int main(int argc, char* argv[])
48 {
49     double start, finish;
50     /*Get number of threads from command line*/

```

```

51     Get_args(argc,argv);           /*Inputs*/
52     points_per_thread = total_points / thread_count;
53
54     /*Set random seed*/
55     srand((unsigned)time(NULL));
56     pthread_t *threads = malloc(thread_count *
sizeof(pthread_t));
57
58     /*Initialize the properties of thread*/
59     pthread_attr_t attr;
60     pthread_attr_init(&attr);
61     GET_TIME(start);
62     int i;
63     for (i = 0; i < thread_count; i++) {
64         pthread_create(&threads[i], &attr, Thread_sum, (void *)
NULL);
65     }
66
67     for (i = 0; i < thread_count; i++) {
68         pthread_join(threads[i], NULL);
69     }
70     GET_TIME(finish);
71     /*free mutex and threads*/
72     pthread_mutex_destroy(&mutex);
73     free(threads);
74
75     printf("Using Monte Carlo Method estimate Pi: %f\n", (4. *
(double)in_circle) / ((double)points_per_thread *
thread_count));
76     printf("The elapsed time is : %f\n",finish-start);
77
78     return 0;
79 }
80
81 /*-----
-----
82  * Function:    Get_args
83  * Purpose:     Get the command line args
84  * In args:     argc, argv
85  * Globals out: total_points, thread_count
86  */
87 void Get_args(int argc, char* argv[]) {
88     if (argc != 3) usage(argv[0]);

```

```

89     total_points = strtol(argv[1], NULL, 10);
90     if (total_points <= 0) Usage(argv[0]);
91     thread_count = strtoll(argv[2], NULL, 10);
92     if (thread_count <= 0 || thread_count > MAX_THREADS)
Usage(argv[0]);
93 } /* Get_args */
94
95 /*-----
-----
96  * Function:  Usage
97  * Purpose:   Print a message explaining how to run the program
98  * In arg:    prog_name
99  */
100 void Usage(char* prog_name) {
101     fprintf(stderr, "usage: %s <total points>
<threads>\n", prog_name);
102     fprintf(stderr, "total points should be evenly divisible by
the number of threads\n");
103     fprintf(stderr, "threads is the number of terms and should
be >= 1\n");
104     exit(1);
105 } /* Usage */
106
107 /*-----
-----
108  * Function:  Thread_sum
109  * Purpose:   Accumulate qualified random numbers
110  * In arg:    None
111  * Globals out: in_circle
112  */
113 void *Thread_sum() {
114     long in_circle_thread = 0;
115     unsigned int rand_state = rand();
116     long i;
117     for (i = 0; i < points_per_thread; i++) {
118         /*random() function does not improve performance
119         * therefore chooses rand_f function
120         * */
121         double x = rand_r(&rand_state) / ((double)RAND_MAX + 1)
* 2.0 - 1.0;
122         double y = rand_r(&rand_state) / ((double)RAND_MAX + 1)
* 2.0 - 1.0;
123         if (x * x + y * y < 1) {

```

```

124         in_circle_thread++;
125     }
126 }
127
128 pthread_mutex_lock(&mutex);
129 in_circle += in_circle_thread;
130 pthread_mutex_unlock(&mutex);
131
132 return NULL;
133 }/*Thread_sum*/

```

2.3.2 Screenshot

```

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./pth_pi 10000000 1
Using Monte Carlo Method estimate Pi: 3.141084
The elapsed time is : 0.122662

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./pth_pi 10000000 2
Using Monte Carlo Method estimate Pi: 3.140288
The elapsed time is : 0.061615

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./pth_pi 10000000 4
Using Monte Carlo Method estimate Pi: 3.140509
The elapsed time is : 0.038465

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./pth_pi 10000000 6
Using Monte Carlo Method estimate Pi: 3.139991
The elapsed time is : 0.026679

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./pth_pi 10000000 8
Using Monte Carlo Method estimate Pi: 3.140134
The elapsed time is : 0.022326

```

```
86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./pth_pi 10000000 10
Using Monte Carlo Method estimate Pi: 3.139588
The elapsed time is : 0.029461

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./pth_pi 10000000 12
Using Monte Carlo Method estimate Pi: 3.140147
The elapsed time is : 0.028828

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./pth_pi 10000000 14
Using Monte Carlo Method estimate Pi: 3.139906
The elapsed time is : 0.025320

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./pth_pi 10000000 16
Using Monte Carlo Method estimate Pi: 3.139962
The elapsed time is : 0.029925
```

III. Lab 4 (35 Marks)

3.1 OpenMP Hello World

3.1.1 Code

```
1 gcc -g -Wall -fopenmp -o omp_hello omp_hello.c
2 ./omp_hello 2
3 ./omp_hello 4
4 ./omp_hello 6
```

omp_hello.c

```
1 #include<stdlib.h>
2 #include<stdio.h>
3 #include<omp.h>
4
5 void Hello(void); //Threads function
6
7 int main(int argc, char* argv[]){
8     /*Get number of threads from command line*/
9     int thread_count = strtol(argv[1], NULL, 10);
10    printf("Max threads: %d \n", omp_get_num_procs());
11
12    #pragma omp parallel num_threads(thread_count)
13    Hello();
14
15    return 0;
16 }
17
18 void Hello(void){
19     int my_rank = omp_get_thread_num();
20     int thread_count = omp_get_num_threads();
21
22     printf("Hello from thread %d of %d\n", my_rank, thread_count);
23 }
```

3.1.2 Screenshot

```
86133@Priska /cygdrive/d/dts202tc/Openmp_1ec/Basic
$ gcc -g -Wall -fopenmp -o omp_hello omp_hello.c

86133@Priska /cygdrive/d/dts202tc/Openmp_1ec/Basic
$ ./omp_hello 2
Max threads: 8
Hello from thread 0 of 2
Hello from thread 1 of 2

86133@Priska /cygdrive/d/dts202tc/Openmp_1ec/Basic
$ ./omp_hello 4
Max threads: 8
Hello from thread 0 of 4
Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 3 of 4

86133@Priska /cygdrive/d/dts202tc/Openmp_1ec/Basic
$ ./omp_hello 6
Max threads: 8
Hello from thread 4 of 6
Hello from thread 3 of 6
Hello from thread 0 of 6
Hello from thread 1 of 6
Hello from thread 2 of 6
Hello from thread 5 of 6
```

It is observed that there is competition between threads to access standard output, so the output does not appear in thread numbering order.

3.2 Estimate π with OpenMP

3.2.1 Code

```
1 gcc -g -Wall -fopenmp -o omp_pi omp_pi.c
2 ./omp_pi 10000000 1
3 ./omp_pi 10000000 2
4 ./omp_pi 10000000 4
5 ./omp_pi 10000000 6
6 ./omp_pi 10000000 8
7 ./omp_pi 10000000 10
8 ./omp_pi 10000000 12
9 ./omp_pi 10000000 14
10 ./omp_pi 10000000 16
```

omp_pi.c

```
1  /*File: omp_pi.c
2   * Purpose: Estimate pi using Monte Carlo Method
3   *          (number in circle)/(total number of tosses) = pi/4
4   *
5   * Compile:      gcc -g -Wall -fopenmp -o omp_pi omp_pi.c
6
7   * Run:          ./omp_pi <total_points><number of threads>
8   *              total_points is the randomly toss total
9   *              should be evenly divisible by the number of
10  threads
11  *
12  *Input:      None
13  *Output: The estimate of pi using Monte Carlo on multiple
14  threads,
15  *          Also elapsed times for the multithreaded and
16  singlethreaded
17  *          computations
18  */
19
20 #include<stdlib.h>
21 #include<stdio.h>
22 #include<time.h>
23 #include<omp.h>
```

```

21
22 #ifndef _TIMER_H_
23 #define _TIMER_H_
24 #include <sys/time.h>
25 #include <stdlib.h>
26 /* The argument now should be a double (not a pointer to a
double) */
27 #define GET_TIME(now) { \
28     struct timeval t; \
29     gettimeofday(&t, NULL); \
30     now = t.tv_sec + t.tv_usec/1000000.0; \
31 }
32 #endif
33
34 const int MAX_THREADS = 1024;
35 void Usage(char prog_name);
36
37 int main(int argc, char** argv){
38     double start, finish;
39     long long in_circle=0;
40     long long total_points;
41     int thread_count;
42
43     /*Get number of threads from command line*/
44     if(argc != 3) Usage(argv[0]);
45     total_points = strtol(argv[1], NULL, 10);
46     if(total_points <=0 ) Usage(argv[0]);
47     thread_count=strtol(argv[2], NULL, 10);
48     if(thread_count<=0 || thread_count > MAX_THREADS)
Usage(argv[0]);
49
50     /*Set random seed*/
51     srand(time(NULL));
52
53     /*thread compute local circle points*/
54     double x,y,distance_point;
55     long long int i;
56
57     GET_TIME(start);
58 #pragma omp parallel for num_threads(thread_count) default(none)
\
59     reduction(+:in_circle) shared(total_points)
private(i,x,y,distance_point)

```

```

60     for( i=0;i<total_points;i++){
61         x=(double)rand()/(double)RAND_MAX;
62         y=(double)rand()/(double)RAND_MAX;
63         distance_point=x*x+y*y;
64         if(distance_point<=1){
65             in_circle++;
66         }
67     }
68     double estimate_pi=(double)in_circle/total_points*4;
69     GET_TIME(finish);
70
71     printf("Using Monte Carlo Method estimate Pi:
72 %lf\n",estimate_pi);
73     printf("The elapsed time: %.8f\n",finish-start);
74     return 0;
75 }
76
77 /*-----
78
79 * Function:  Usage
80 * Purpose:   Print a message explaining how to run the program
81 * In arg:    prog_name
82 */
83 void Usage(char* prog_name) {
84     fprintf(stderr, "usage: %s <total points>
85 <threads>\n",prog_name);
86     fprintf(stderr,"total points should be evenly divisible by
87 the number of threads\n");
88     fprintf(stderr,"threads is the number of terms and should be
89 >= 1\n");
90     exit(1);
91 } /* Usage */

```

3.2.2 Screenshot

```

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ gcc -g -Wall -fopenmp -o omp_pi omp_pi.c

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./omp_pi 10000000 1
Using Monte Carlo Method estimate Pi: 3.141423
The elapsed time: 0.09397507

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./omp_pi 10000000 2
Using Monte Carlo Method estimate Pi: 3.140911
The elapsed time: 0.04925394

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./omp_pi 10000000 4
Using Monte Carlo Method estimate Pi: 3.141111
The elapsed time: 0.03069091

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./omp_pi 10000000 6
Using Monte Carlo Method estimate Pi: 3.142101
The elapsed time: 0.02260590

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./omp_pi 10000000 8
Using Monte Carlo Method estimate Pi: 3.142596
The elapsed time: 0.01728702

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./omp_pi 10000000 10
Using Monte Carlo Method estimate Pi: 3.141840
The elapsed time: 0.02691722

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./omp_pi 10000000 12
Using Monte Carlo Method estimate Pi: 3.142404
The elapsed time: 0.02273989

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./omp_pi 10000000 14
Using Monte Carlo Method estimate Pi: 3.143034
The elapsed time: 0.02100992

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./omp_pi 10000000 16
Using Monte Carlo Method estimate Pi: 3.143695
The elapsed time: 0.01811695

```

3.3 Comparison

3.3.1 Comparison under different threads

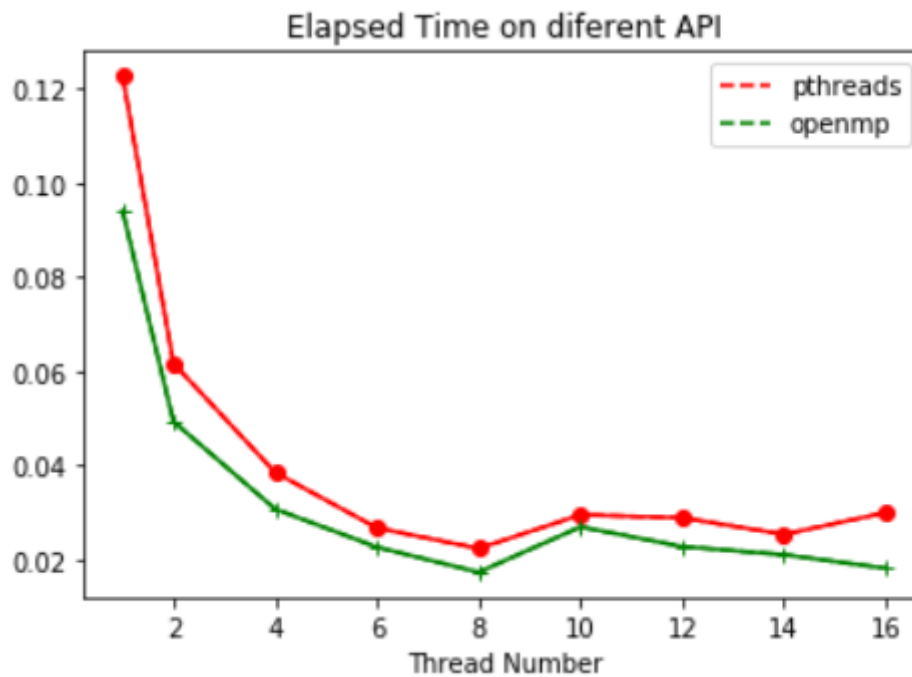
3.3.1.1 Elapsed Time

The following is a table of time under different threads using two different API, Pthreads and OpenMP

	1 THREAD	2 THREADS	4 THREADS	6 THREADS	8 THREADS
Pthreads	0.1226 s	0.06161 s	0.03846 s	0.02667 s	0.02232 s
OpenMP	0.09397 s	0.04925 s	0.03069 s	0.02260 s	0.01728 s

	10 THREADS	12 THREADS	14 THREADS	16 THREADS
Pthreads	0.02946 s	0.02882 s	0.02532 s	0.02992 s
OpenMP	0.02691 s	0.02273 s	0.02100 s	0.01811 s

The following figure compares the time consumption of `Pthreads` and `OpenMP` through a line chart.

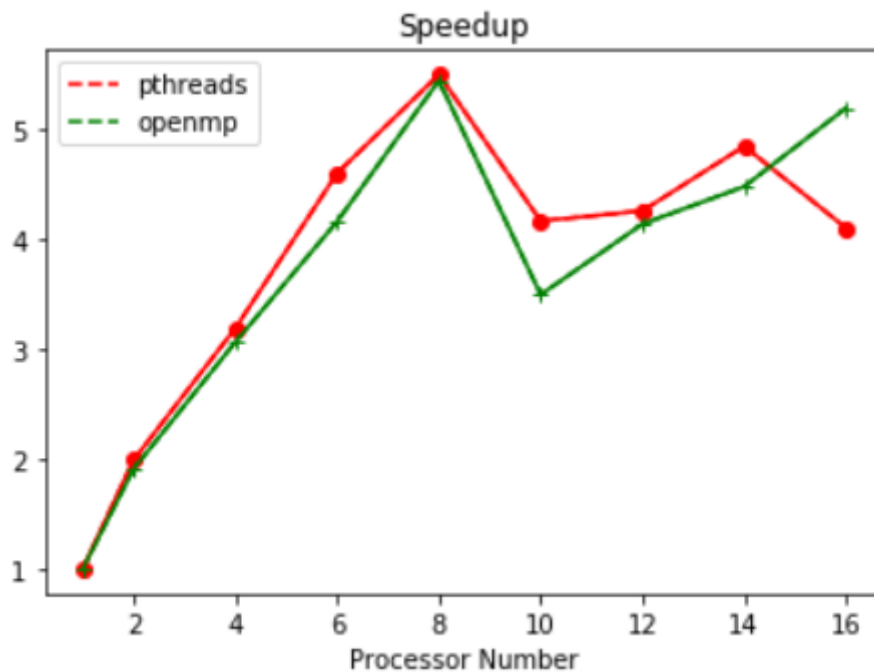


The author finds that when using `OpenMP`, there is less time for each process than when using `Pthreads`. Next, look at speedup and efficiency.

3.3.1.2 Speedup

	1 THREAD	2 THREADS	4 THREADS	6 THREADS	8 THREADS
Pthreads	1	1.9899	3.1877	4.5969	5.4928
OpenMP	1	1.9080	3.0619	4.1579	5.4380

	10 THREADS	12 THREADS	14 THREADS	16 THREADS
Pthreads	4.1615	4.2539	4.8420	4.0975
OpenMP	3.4920	4.1341	4.4747	5.1888



When running parallel programs with 1-8 threads, the speedup of the two is almost the same, and it is found that before 8 threads, speedup is often close to linear speedup.

In addition, there are some tasks that cannot be parallelized in the program. According to **Amdahl's Law**, there are parts with a ratio of r in serial programs that cannot be parallelized, and the speedup that can be achieved is close to $\frac{1}{r}$ [1]

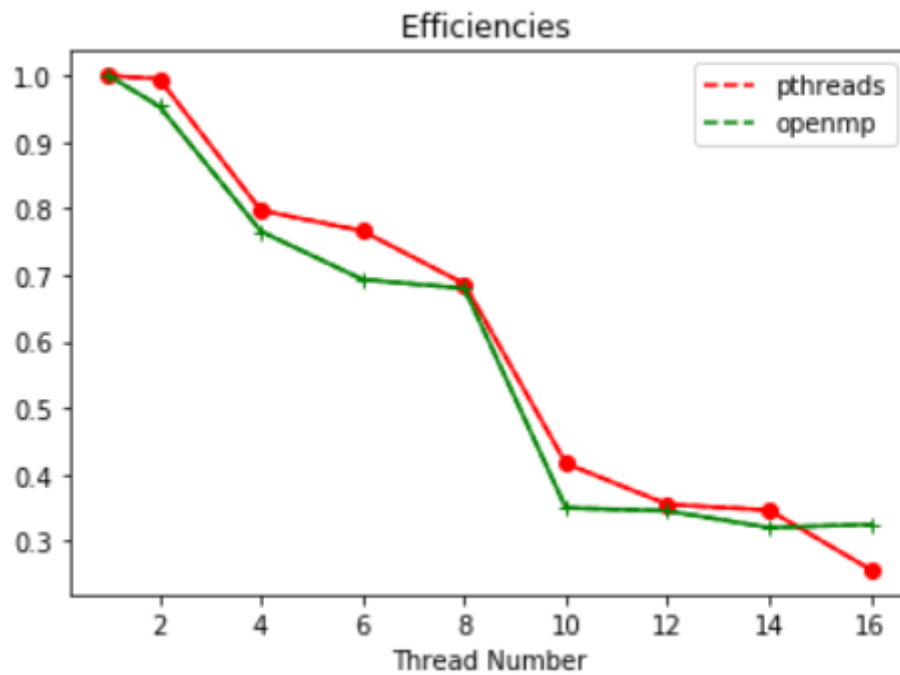
Both Pthreads and openmp reach the maximum speedup of 5.5 at 8 threads, so the non-parallelization task of both is 18%.

When more than 8 threads run a parallelized program, author will find that the speedup of the program decreases or fluctuates in an area. This phenomenon is that when the number of threads increases, the busy waiting time between threads will increase, resulting in a significant increase in overhead time, and the speedup will fluctuate.

3.3.1.3 Efficiency

	1 THREAD	2 THREADS	4 THREADS	6 THREADS	8 THREADS
Pthreads	1	0.9949	0.7969	0.7661	0.6866
OpenMP	1	0.9540	0.7654	0.6797	0.6797

	10 THREADS	12 THREADS	14 THREADS	16 THREADS
Pthreads	0.4161	0.3544	0.3458	0.2560
OpenMP	0.3492	0.3445	0.3196	0.3243



$$E = \frac{\text{speedup}}{p}, p \text{ is pthreads number};$$

As the number of threads increases, the efficiency decreases gradually. This is because the more threads, the higher the percentage of overhead time.

3.3.2 Comparison under different problem size

When problem size = 5000000

```
86133@Priska /cygdrive/d/dts202tc/pthreads_lab/lab2
$ gcc -g -Wall -o pth_pi pth_pi.c -lm -lpthread

86133@Priska /cygdrive/d/dts202tc/pthreads_lab/lab2
$ ./pth_pi 5000000 1
Using Monte Carlo Method estimate Pi: 3.140288
The elapsed time is : 0.061320

86133@Priska /cygdrive/d/dts202tc/pthreads_lab/lab2
$ ./pth_pi 5000000 2
Using Monte Carlo Method estimate Pi: 3.140509
The elapsed time is : 0.031455

86133@Priska /cygdrive/d/dts202tc/pthreads_lab/lab2
$ ./pth_pi 5000000 4
Using Monte Carlo Method estimate Pi: 3.140134
The elapsed time is : 0.020605

86133@Priska /cygdrive/d/dts202tc/pthreads_lab/lab2
$ ./pth_pi 5000000 8
Using Monte Carlo Method estimate Pi: 3.139962
The elapsed time is : 0.011450

86133@Priska /cygdrive/d/dts202tc/pthreads_lab/lab2
$ gcc -g -Wall -fopenmp -o omp_pi omp_pi.c

86133@Priska /cygdrive/d/dts202tc/pthreads_lab/lab2
$ ./omp_pi 5000000 1
Using Monte Carlo Method estimate Pi: 3.140450
The elapsed time: 0.04595280

86133@Priska /cygdrive/d/dts202tc/pthreads_lab/lab2
$ ./omp_pi 5000000 2
Using Monte Carlo Method estimate Pi: 3.141502
The elapsed time: 0.02434182

86133@Priska /cygdrive/d/dts202tc/pthreads_lab/lab2
$ ./omp_pi 5000000 4
Using Monte Carlo Method estimate Pi: 3.141934
The elapsed time: 0.01615787

86133@Priska /cygdrive/d/dts202tc/pthreads_lab/lab2
$ ./omp_pi 5000000 8
Using Monte Carlo Method estimate Pi: 3.143777
The elapsed time: 0.00961304
```

	1 THREAD	2 THREADS	4 THREADS	8 THREADS
Pthread	0.06132	0.031455	0.02060	0.01145
OpenMP	0.04595	0.02434	0.01615	0.00961

When problem size = 20000000

```

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./pth_pi 20000000 1
Using Monte Carlo Method estimate Pi: 3.141517
The elapsed time is : 0.245479

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./pth_pi 20000000 2
Using Monte Carlo Method estimate Pi: 3.141084
The elapsed time is : 0.123934

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./pth_pi 20000000 4
Using Monte Carlo Method estimate Pi: 3.140288
The elapsed time is : 0.071555

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./pth_pi 20000000 8
Using Monte Carlo Method estimate Pi: 3.140509
The elapsed time is : 0.042611

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./omp_pi 20000000 1
Using Monte Carlo Method estimate Pi: 3.141707
The elapsed time: 0.18290591

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./omp_pi 20000000 2
Using Monte Carlo Method estimate Pi: 3.141690
The elapsed time: 0.09538984

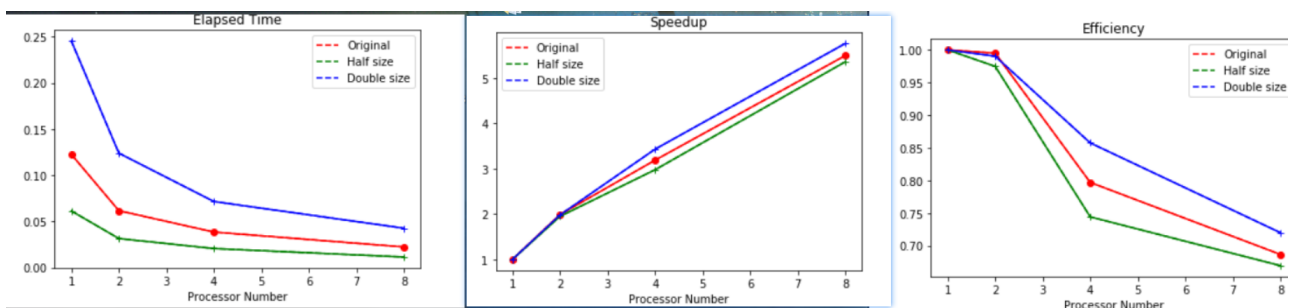
86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./omp_pi 20000000 4
Using Monte Carlo Method estimate Pi: 3.141600
The elapsed time: 0.05485678

86133@Priska /cygdrive/d/dts202tc/threads_lab/lab2
$ ./omp_pi 20000000 8
Using Monte Carlo Method estimate Pi: 3.141113
The elapsed time: 0.03380394

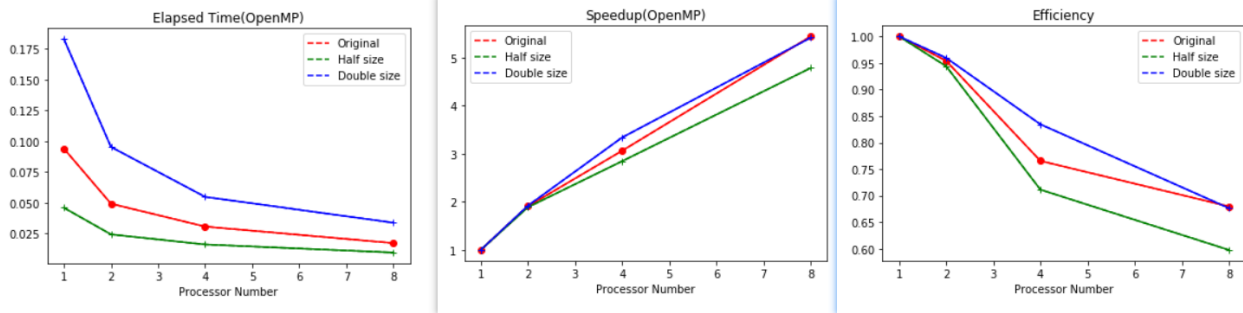
```

	1 THREAD	2 THREADS	4 THREADS	8 THREADS
Pthread	0.2454	0.1239	0.07155	0.04261
OpenMP	0.1829	0.0953	0.0548	0.03380

Pthread:



OpenMP:



The values of Elapsed time, Speedup and Efficiency depend not only on the number of threads, but also on the scale of the problem. When the scale of the problem grows, speedup and efficiency increase at the same time.

According to:

$$T_{Parallel} = \frac{T_{Serial}}{p} + T_{overhead}$$

As the scale of the problem increases, threads have more tasks to perform, and the relative time for coordinating work between threads is reduced. overhead time grows more slowly than serial time, consequently both Speedup and efficiency will increase.

And according to Gustafson's Law, as the scale of the problem increases, the proportion of the nonparallelizable part of the program also decreases [2].

From half the size to the original problem size and then to double the size, the maximum speedup is increasing, so according to the formula, the non-parallelization part is reduced.

- Reference List

[1] Krishnaprasad, S., 2001. Uses and abuses of Amdahl's law. Journal of Computing Sciences in colleges, 17(2), pp.288-293.

[2] McCool, Michael D.; Robison, Arch D.; Reinders, James (2012). "2.5 Performance Theory". Structured Parallel Programming: Patterns for Efficient Computation. Elsevier. pp. 61–62. ISBN 978-0-12-415993-8